

TinyMAC: From 2×2 8-bit Mac Array to 4×4 Micro-Accelerator

Huancheng Yang Jiangyue Chu

NYU Tandon School of Engineering

CogniChip Hackathon — Feb 20, 2026



Repository

GitHub: <https://github.com/Torzirael36/Cognichip-Hackson.git>

Outline

- 1 Introduction
- 2 Architecture and RTL
- 3 Using the Cognichip Platform
- 4 Verification and Simulation Evidence

Problem & Motivation

- ▶ Multiply-Accumulate (MAC) is the core operation behind matrix multiplication and convolution:

$$\text{acc} \leftarrow \text{acc} + (a \times b)$$

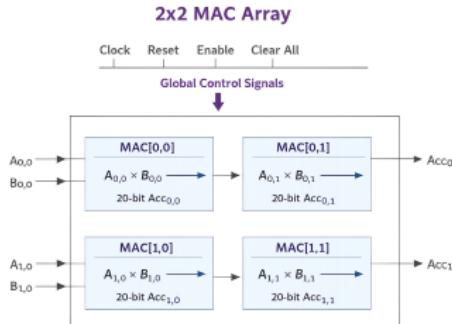
- ▶ Modern AI and DSP pipelines perform massive numbers of MAC operations; small compute blocks are the building units.
- ▶ **Goal:** build a compact, synthesizable **2×2 8-bit MAC array** with simple control and strong simulation evidence.

What we submit

RTL (SystemVerilog) + simulation testbench evidence
(waveforms/logs) + final slides + demo video.

What We Built: TinyMAC 2×2

- ▶ A hierarchical, modular design with two layers:
 - ▶ **mac_unit**: single MAC with 8-bit inputs and a wide accumulator
 - ▶ **mac_array_2x2**: 4 MACs in parallel with global control signals
- ▶ Parallel compute: **4 MAC operations per cycle** (when enabled)
- ▶ Simple control interface: enable (pause/resume) and clear_all (reset accumulators)



What We Built: TinyMAC 4×4 Micro-Accelerator

- ▶ **Extended (4×4 micro-accelerator):**
 - ▶ **Throughput:** 16 MACs in parallel **(16 MACs/cycle)**
 - ▶ **A/B register buffers:** 4×4 each (A_buf, B_buf)
 - ▶ **N-cycle accumulation:** matrix-multiply style with k loop
($k=0..3$)
 - ▶ **Controller FSM:** IDLE/LOAD/CLEAR/SETUP/COMPUTE/DONE
 - ▶ **Readout:** register-mapped interface out_addr / out_rdata
 - ▶ **Verification:** self-checking SV testbench + golden model + waveform/log evidence

MAC Semantics and Timing

- ▶ Operation:

$$\text{accumulator} \leftarrow \text{accumulator} + (a \times b)$$

- ▶ **Synchronous accumulation:** updates occur on the rising edge of clock.
- ▶ **Reset/Clear behavior:**
 - ▶ reset or clear forces accumulator to zero
 - ▶ enable gates accumulation (no update when disabled)
- ▶ Product width: 8-bit \times 8-bit \Rightarrow 16-bit product.

Design choice

We use a wider accumulator to reduce overflow risk during repeated accumulation.

Accumulator Width Choice ($\text{int8} \times \text{int8} \rightarrow \text{int32}$)

- ▶ Default parameters (4×4 accelerator RTL):
 - ▶ DATA_W = 8 (signed int8 inputs, range: $[-128, 127]$)
 - ▶ ACC_W = 32 (signed accumulator)

- ▶ Single-product bound (signed int8):

$$\max |a \cdot b| = 128 \times 128 = 16384 \quad (\text{fits in 16 bits})$$

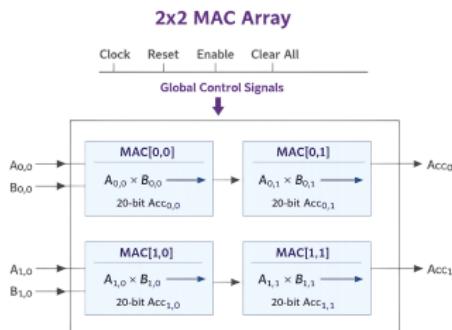
- ▶ For 4×4 matmul, each output accumulates $N = 4$ terms:

$$\max |C[i][j]| \leq 4 \times 16384 = 65536 \quad (\approx 17 \text{ bits magnitude})$$

- ▶ Using a 32-bit accumulator provides robust headroom for sign, corner cases, and future scaling (e.g., larger N or deeper accumulation).

Module 2: mac_array_2x2

- ▶ A 2×2 grid of 4 independent MAC units:
 - ▶ MAC[0,0], MAC[0,1], MAC[1,0], MAC[1,1]
- ▶ Shared global signals:
 - ▶ clock, reset, enable, clear_all
- ▶ Independent channels:
 - ▶ each MAC has its own a_{ij} , b_{ij} inputs and acc_{ij} output



4×4 Accelerator Micro-architecture

Compute Target (Matrix Multiply Kernel)

$$C[i][j] = \sum_{k=0}^3 A[i][k] \cdot B[k][j]$$

Architecture Highlights

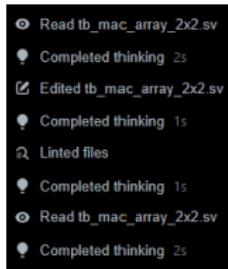
- ▶ **Parallelism:** 16 MACs per cycle (all i, j in parallel)
- ▶ **N-cycle accumulation:** $N = 4$ cycles over k
- ▶ **Per-cycle dataflow:**
 - ▶ fixed k_index
 - ▶ broadcast $A[:, k]$ and $B[k, :]$ into array
- ▶ **Internal state:**
 - ▶ A_buf [4] [4]
 - ▶ B_buf [4] [4]
 - ▶ C_acc [4] [4]

How We Used Cognichip

- ▶ Generated and iterated on SystemVerilog RTL modules using the Cognichip workflow.
- ▶ Created/updated simulation targets (RTL + testbench + run configuration).
- ▶ Ran simulation to produce logs and waveform traces for correctness evidence.

How We Used Cognichip for RTL Development

- ▶ We used Cognichip as an AI-assisted RTL co-design platform.
- ▶ Started with a high-level architectural description:
 - ▶ 2x2 MAC array
 - ▶ 8-bit inputs
 - ▶ 20-bit accumulator
 - ▶ Global control signals
- ▶ Cognichip generated synthesizable SystemVerilog automatically.
- ▶ Built hierarchical modules:
 - ▶ mac_unit
 - ▶ mac_array_2x2



Prompt Engineering: Guiding the Design

- ▶ Instead of vague requests, we provided:
 - ▶ Explicit data widths
 - ▶ Accumulator overflow constraints
 - ▶ Clear/reset semantics
 - ▶ Hierarchical module structure
- ▶ Structured design approach:
 - ▶ Step 1: Define mac_unit
 - ▶ Step 2: Instantiate 2x2 grid
 - ▶ Step 3: Add control signals
- ▶ We treated Cognichip like a hardware co-engineer.

Completed thinking 2s
Great! Let's build a simple 2x2 8-bit MAC (Multiply-Accumulate) Array together.

Design approach:

- We'll create a 2x2 grid of MAC units (4 total)
- Each MAC performs: accumulator = (accumulator * input.a) + input.b
- 8-bit inputs, with wider accumulator to prevent overflow (we'll use 20 bits to safely hold accumulated products)
- Simple interface: clock, reset, enable, and separate inputs for each MAC position
- Each MAC can be cleared independently or all at once

I'll create a hierarchical design with:

1. A single MAC unit module
2. A top-level module instantiating 4 MAC units in a 2x2 array

Let me write the code:

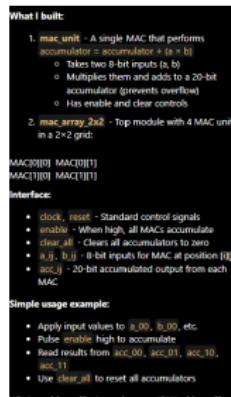
Created mac_array_2x2.sv

Completed thinking 2s

Listed files

Iterative Refinement and Self-Verification

- ▶ After initial generation:
 - ▶ Ran lint checks
 - ▶ Requested improvements
 - ▶ Verified coding style
- ▶ Cognichip automatically:
 - ▶ Corrected structural issues
 - ▶ Ensured synthesizable constructs
 - ▶ Maintained clean hierarchical separation
- ▶ Achieved zero lint warnings.



Why Cognichip Is Powerful

- ▶ Rapid RTL generation:
 - ▶ Full 2x2 MAC array built in seconds
 - ▶ Clean, parameterized, synthesizable code
 - ▶ Supports hierarchical hardware abstraction
 - ▶ Enables faster design exploration

Engineering Impact

Cognichip significantly reduces development time while maintaining hardware-level correctness and modularity.

From idea to working RTL in minutes.

Verification Methodology (4×4)

Golden Reference Model

- ▶ SystemVerilog matrix-multiply reference
- ▶ Signed int8 inputs → int32 accumulation

Directed Tests

- ▶ all zeros
- ▶ identity × random ($A = I \Rightarrow C = B$)
- ▶ max positive / max negative
- ▶ mixed signs

Random Regression

- ▶ 5–20 random matrix pairs per run
- ▶ full 16-element compare

Verification Checks and Evidence

Correctness Checks

- ▶ done latency: CLEAR + SETUP + N compute cycles
- ▶ Per-element comparison
 - ▶ All 16 outputs verified
 - ▶ via out_addr / out_rdata

Artifacts

- ▶ Waveform inspection
- ▶ PASS / FAIL summary logs
- ▶ Automatic mismatch detection

Simulation Results: 2×2 8-bit Mac Array Waveform Evidence



Test Case Walkthrough: Multiply-Accumulate Behavior

► Step 1: $3 \times 4 = 12$

- Inputs: $a_{00} = 3$, $b_{00} = 4$, enable = 1
- Product: $3 \times 4 = 12 = 0x000C$
- Accumulator update:

$$acc_{00} : 0x0000 \rightarrow 0x000C$$

- acc_00 matches expected_00

► Step 2: $5 \times 6 = 30$

- $5 \times 6 = 30 = 0x001E$
- Accumulation:

$$0x000C + 0x001E = 0x002A \text{ (42)}$$

- Waveform shows step-wise increase to 0x002A

► Step 3: $2 \times 3 = 6$

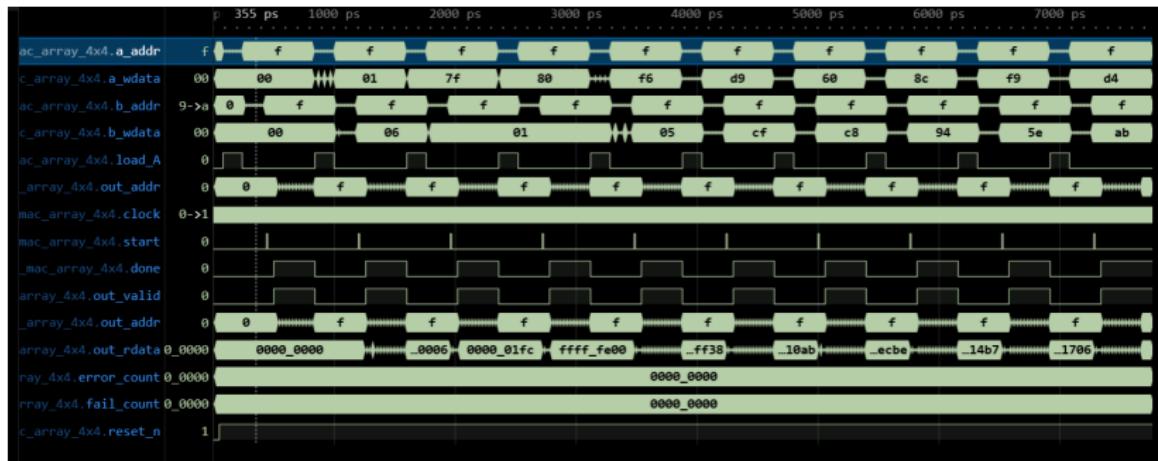
- $2 \times 3 = 6$
- Accumulation:

$$0x002A + 0x0006 = 0x0030 \text{ (48)}$$

Simulation Results: Log Evidence

- ▶ The cognichip runs the testbenches.
- ▶ All directed tests passed (9/9).
- ▶ The reset,enable signals are all tested.
- ▶ Logs record key checks and final PASS/FAIL summary.

Simulation Results: 4×4 Accelerator Waveform Evidence



4×4 Example: Identity × Random (Sanity Check)

- ▶ Test setup:
 - ▶ Set $A = I$ (4×4 identity), random B
 - ▶ Expected result: $C = A \cdot B = B$
- ▶ Test flow (self-checking):
 - ▶ TB loads A_buf and B_buf via `load_A/load_B + address/data`
 - ▶ Assert `start`; controller runs $k=0..3$ accumulation
 - ▶ Read back all 16 outputs via `out_addr/out_rdata`
 - ▶ Compare against golden model (SV reference `matmul`)
- ▶ Takeaway: validates end-to-end datapath + control + readout interface

Simulation Results: 4×4 Log Evidence

- ▶ Directed tests: **5/5 passed**
- ▶ Random regression: **5/5 passed**
- ▶ Total mismatches: **0**
- ▶ Testbench structure: Tests 1–5 directed, Tests 6–10 random
- ▶ Logs include per-element checks and final PASS/FAIL summary.

```
=====
TEST SUMMARY
=====
Total Tests: 10
Passed:    10
Failed:    0
Errors:    0
=====
TEST PASSED
```

Challenges, Lessons, and Future Directions

Technical Challenges

- ▶ N -cycle accumulation with correct k-loop control
- ▶ Off-by-one / skip-k FSM sequencing bugs
- ▶ Clean load–compute–readout interface

Key Lessons

- ▶ Verification must validate microarchitecture, not just arithmetic
- ▶ Self-checking TB exposed real control-path misalignment
- ▶ Modular design (MAC → array → controller) enables scaling

Future Directions

- ▶ Scale to 8×8 and larger arrays
- ▶ Add pipelining / retiming for higher clock frequency
- ▶ Explore systolic dataflow variants
- ▶ FPGA demo and PPA measurement